Niva job finding verification

Independent retest of the original audit and repair claims

Overall verdict: verified with remaining issues. The live release contains the reviewed discovery repairs. Strict hybrid filtering now exposes the two intended roles, the false collaboration skill gap is removed, and four salary omissions are corrected. Description semantics remain incomplete and a freshness summary is misleading. The evidence does not establish that every issue is resolved or that all displayed vacancies are open and eligible.

Audit verdict	Count
Passed	14
Partially passed	14
Failed	1
Not reproducible	0
Blocked / not tested	4

Counts cover 33 audit and feature rows: original B01–B04 and S01–S03, plus F01–F26 labels assigned by the fix report to the audit's unnumbered scorecard. These labels are preserved for traceability. One new low-severity issue N01 is counted separately as Failed. F17 is an unimplemented optional recommendation, not a release regression.

Release decision

The checks support the specific repairs already deployed, with remaining issues. They do not support an unconditional production-readiness sign-off. S02 still needs accountable hosted retention, alerting, identity-isolation and restoration evidence. Live write paths remain untested under the user's explicit instruction. No critical discovery or application blocker was demonstrated in the inspected public cohort; no successful application submission is claimed.

Tested version

Item	Recorded evidence
Environment	https://getnivajobs.com/app/jobs · production read-only; local isolated Node24 SQLite and D1-contract tests
Date and browser	26 October 2026; public collection17:04–17:05 UTC; manual UI 17:07–17:17 UTC; Windows IAB Chromium engine; exact embedded browser build not exposed
Viewports	Desktop1440×1000; mobile 390×844 emulation; local Chromium desktop and mobile in both themes; no physical phone
Deployed release	0.2.1 / Sites41; build-info built16:48:38 UTC; reviewed GitHub 55b21e270dab0ec25cd8d908ebdf600425d45cbb
Deployment provenance	Sites source ddc96b3a18b9e5861a0b48208b1dff6ea66c12e0; mapped-reviewed-source true; dirty false
Local checkout	8550b51a7247039caa44424c3f1c5c4b90284458; src/server/package files equal reviewed 55b21e2; later local commit is documentation reconciliation

Method and authorization

I read all 28 audit pages and all 15 fix-report pages, including the supplied screenshots, scorecard, tested scenarios and bug register. The fix report describes an earlier local repair state; its statement that changes were uncommitted and unpublished is historical. Current deployment provenance and observed live behavior establish that the reviewed fixes are now deployed. Previous test claims were not accepted as new passes.

The original persona was retained: five years in QA, mid/senior individual contributor, Python, TypeScript, Playwright, Selenium, SQL, APIs and GitHub Actions; Tel Aviv commute, onsite/hybrid/remote; Israel-only authorization and no US citizenship; ILS25000–35000 gross monthly desired pay. Degree, Java, leadership, gaming and AI-agent depth remain unconfirmed. The exact seven skills were used in the independent isolated assessment. Existing browser regression fixtures use a smaller four-skill synthetic profile, so their scores are not substituted for persona calibration.

The user authorized isolated local data and requested live writes remain untested. No live profile, résumé, preference, saved-job, onboarding-account or application write was performed. The live Apply redirect is intentionally excluded because its GET handler inserts an apply-click event. Original employer pages were opened directly instead, and redirect behavior was tested locally. Ordinary feed/detail GET requests can update shared cache and provider quota metadata internally; they were not forced refresh or load tests.

Evidence register

E1 Fresh public GET data 17:04:31–17:05:08 UTC; 20 details, 14 official provider APIs, 20 original HTTP paths, seven scenarios and two-page traversal. public.json.

E2 Independent live IAB UI 17:07–17:17 UTC; desktop 1440×1000 and mobile 390×844. Counts, history, chips, map, sorting, empty recovery, keyboard and freshness inspected. Screenshots in verify41.

E3 Fourteen original source pages visibly inspected 17:08:44–17:09:54 UTC. browser-sources.json. Six blank application forms observed; no field filled or submit clicked.

E4 Exact audit persona assessed against 20 freshly fetched jobs in isolation. assessments.json and assess.ts; 28 assertions passed. Not a live profile update or empirical fit calibration.

E5 Fresh full local execution: frontend 59/59; server 153/153; Chromium 218/218 across desktop/mobile; Worker 1/1; provenance 1/1; build and three type checks passed. No failed, skipped or flaky browser tests.

E6 Original supplied 28-page audit and 15-page fix report read in full, including scores, screenshots and ledgers. Original images are explicitly labelled.

E7 Unperformed dependencies: hosted identity/retention/alerts, live profile/CV/save writes, live apply tracking redirect, real AI, field performance and physical-device/screen-reader validation.

What the sample can establish

All20 original unique results were re-read from fresh public Niva details; fourteen were independently compared with current official Greenhouse/Lever payloads and fourteen original pages were inspected visibly. The original twenty-record purposive cohort is kept separate from inventory changes and from the signed-in owner’s display. Neither a reachable URL nor an Apply control proves the vacancy accepts new applications. No vacancy-wide closure rate, hiring outcome, employer acceptance or market-wide relevance estimate is inferred.

Search scenarios and result quality

Scenario	Original	Fix report catalog	Fresh public	Interpretation
S1 QA automation	4	4	4	Exact category search
S2 QA all categories	7	8	8	One additional LQA posting
S3 QA Tel Aviv	6	6	6	Same district count
S4 remote QA IL including unknown	1	2	2	Country unknowns permitted
S4b strict known IL	0	0	0	No invented eligibility
S5 senior plus ILS30000/month	0	0	0	Honest restrictive result
S6 remote Engineer	338	373	396	Changing inventory; broad fallback is noisy

All seven fresh public requests returned 200. The original URLs were replayed with scope=all and limit 50. During later signed-in manual testing inventory moved to 397 and 49 cards were initially visible, then 99 after load more; all 99 links were unique. Confirmed owner eligibility can affect visible cards, so this is not treated as the same anonymous denominator. The public two-page traversal returned 100 unique IDs with no repeats.

Metric and definition	Original audit	Independent retest
Strong / partial / low role fit in fixed 20	1 /5 /14	1 /5 /14; unchanged independent judgment
Strong or partial role fit	6/20 30%	6/20 30%; not confirmed eligible
Focused QA role-family relevance	6/7 85.7%	6/7 in original QA subset; junior R05 remains low
Known work style	14/20 70%	17/20 85%; R03 R04 unknown and R10 conflict
Structured pay	1/20 5%	3/20 15%; plus R11/R12 pay present with unknown period
Literal amp/nbsp description tokens	13/20 65%	0/20; semantic heading loss still reproduced
Explicit candidate conflicts for S01	Conditions not surfaced reliably	4/4 original citizenship/local-rights cases blocked in isolated exact persona
Closed or broken visibly observed	0/15 original pages	0/14 current pages; subsets differ; not evidence all 20 open
Blank application forms visible	6 original source forms	6 current forms R01 R03 R04 R05 R11 R12; no submission
Distinct canonical vacancy paths	20/20; semantic duplicates0/20	20/20; no conflicting same vacancy found; whole-market rate unmeasured
Numeric score presence	0/20 withheld	7/20 isolated persona; accuracy N/T; no live owner score exported

Publisher integration supplement: the live JSearch view showed 44 in-site cards, including 30 LinkedIn and 14 Indeed labels. It disclosed stale-source status and limited daily coverage. These44 cards are outside the original 20-result quality cohort and were not individually validated. No quota exhaustion or paid call was triggered.

Twenty result review ledger

Independent role fit concerns the original persona. Eligibility is separate: confirmation means unresolved information, not confirmed eligibility. Scores below are isolated exact-persona outputs only. B means a confirmed blocker; Q means confirmation required. Source page V means visibly matched title/content, H means HTTP path only. All20 HTTP paths returned 200; official APIs matched14 IDs/titles.

ID	Role	Role assessment and unresolved gap	Style	Eligibility / score	Page
R01	Canopy QA & Automation Engineer	Strong Strong role fit; 2+ years and API testing. Manual QA and AI-assisted tooling still need confirmation.	hybrid	Q / 73	V
R02	Transmit Security QA & Automation Lead	Partial Partial: 5+ years automation and familiar frameworks; leadership and data-science test depth unresolved.	hybrid	Q / N/S	H
R03	Unframe QA Engineer	Partial Partial: QA/API fit; manual emphasis and AI testing experience unresolved; collaboration gap removed.	unknown	Q / 43	V
R04	Unframe AI-First QA Automation Engineer	Partial Partial: Playwright and experience align; proven agent harness and orchestration expertise unresolved.	unknown	Q / 75	V
R05	WeDev Junior QA Tester	Low Low: junior manual tester below target seniority; JavaScript is required and remains distinct from Java.	onsite	Q / N/S	V
R06	Playson Manual QA Engineer (Game Client)	Partial Partial: manual/API/browser QA; gaming experience and Europe employment inclusion unresolved.	remote	Q / N/S	H
R07	Veeva QA Automation Engineer - CDH	Partial Partial: QA automation; Java, data-platform depth and qualification requirements unconfirmed.	onsite	Q / N/S	V
R08	Datadog Enterprise Sales Engineer - UK/Ireland	Low Low: enterprise sales engineering outside QA; UK/Ireland regional remote.	remote	B / 25	H
R09	Datadog Senior Sales Engineer - Key Accounts	Low Low: sales engineering outside QA; 7+ years; US regional remote. Estimate not an ILS salary match.	remote	B / 18	V
R10	Datadog Senior Product Security Engineer - AI Agents	Low Low: specialized security and AI agents; US locations; source hybrid/remote conflict remains unknown.	unknown	B / N/S	H
R11	JFrog Senior Solution Engineer, Presales - Public Sector (US Remote)	Low Low: public-sector presales/DevOps; explicit US citizenship conflicts with persona.	remote	B / 25	V
R12	JFrog Regional Vice President, Strategic Solution Engineering (Central)	Low Low: regional vice-president management; 10+ years and management depth beyond persona.	remote	B / N/S	V
R13	Wiz Solutions Engineer	Low Low: cloud/security presales; Japan work rights and no sponsorship conflict with persona.	remote	B / N/S	V

ID	Role	Role assessment and unresolved gap	Style	Eligibility / score	Page
R14	Wiz Security Engineer - Product	Low Low: specialized security 7+ years; UK work rights and no sponsorship conflict with persona.	remote	B / N/S	V
R15	Wiz Principal Solutions Engineer	Low Low: principal cloud solutions; Australia rights and no sponsorship conflict with persona.	remote	B / N/S	V
R16	Fieldguide Staff Software Engineer, App Platform	Low Low: staff app-platform development, USA; not QA; USD pay remains not employer verified.	remote	Q / N/S	V
R17	eMerchantPay Senior AI Engineer	Low Low: senior AI/ML engineering in Bulgaria; domain and country mismatch unresolved.	remote	Q / 30	H
R18	Canonical Junior Linux Kernel Engineer - Ubuntu	Low Low: junior kernel/C/Rust rather than QA; Anywhere board label does not prove IL hiring.	remote	Q / N/S	V
R19	Canonical Senior/Staff/Principal Engineer	Low Low: senior systems engineering; deep coding/open-source and travel terms unresolved.	remote	Q / N/S	H
R20	Canonical Cloud Field Engineer	Low Low: cloud consulting/infrastructure rather than QA; employment terms and board gate unresolved.	remote	Q / N/S	V

N/S means score withheld, not zero. Country association inferred from a location is weaker than an employer employment restriction. Remote Anywhere does not prove Israel employment. Four S01 conflicts are explicit authorization/citizenship conditions; other blocked outputs may rely on regional applicant-country restrictions and are not counted as the same four explicit cases.

Finding by finding verification

Expected and actual results are scoped to the tests performed. Partially passed includes successful isolated tests with a live acceptance criterion still untested. Recommendations are evaluated as requested functionality, not retroactively recast as confirmed bugs. Evidence IDs refer to the register.

Audit ID	Original issue	Claimed fix	Test performed	Expected result	Actual result	Verdict	Evidence
B01	Explicit hybrid and onsite arrangements omitted	Recover explicit source styles and retain conflicts	Live strict hybrid search; list/map; fresh 20 records; local conflict controls	Canopy and Transmit retained; onsite recovered; true unknowns honest	Strict hybrid 2 results; styles hybrid/hybrid/onsite/onsite for R01/02/05/07; R10 unknown	Passed	E1 E2 E4 E5
B02	Collaboration becomes Product Management skill gap	Separate required preferred and collaboration evidence	R03 live description versus match component; fresh required facts; positive and negation tests	No collaboration gap; genuine required skill retained	R03 mentions managers but has no Product Management requirement or gap; controls pass	Passed	E1 E2 E5
B03	Entities and flattened description structure	Decode safely and render sections and lists	20 fresh descriptions; Canopy employer h2 versus live tags; mobile; unsafe fixtures	No entity tokens; original section order and semantic headings retained	0/20 amp/nbsp tokens; 27 Canopy list items; Your Mission and opening h2 rendered as paragraphs	Partially passed	E1 E2 E3 E5 Fig 2
B04	Salary absent despite four monetary ranges	Keep currency period origin or unstructured label	Fresh R09 to R12 facts; R09 R11 R12 originals; unit filters in local tests	Four pay disclosures; no inferred JFrog period or currency conversion	Datadog 149000–198000 and 192000–240000 USD/year estimates; JFrog ranges present with unknown period	Passed	E1 E3 E5 Fig 3
S01	Citizenship and local rights omitted suspected	Surface conditions and compare confirmed candidate criteria	Fresh R11 R13 to R15; employer pages; exact audit persona locally; unknown controls	Confirmed conflict blocked; absent candidate rights unresolved	US citizenship and JP GB AU rights retained; four persona conflicts blocked; unknown controls pass	Passed	E1 E3 E4 E5
S02	Overdue maintenance needs operator verification suspected	Local cleanup safeguards; hosted operations remain prerequisite	Reviewed operational claim; reran isolated lifecycle tests only	Two hosted cleanup cycles alerts operator and restore evidence	No authorized hosted operational exercise; no retention breach established; exact dependency remains	Blocked / not tested	E5 E7

Audit ID	Original issue	Claimed fix	Test performed	Expected result	Actual result	Verdict	Evidence
S03	Optional skills years JavaScript and source facts suspected	Respect optional sections exact terms and provenance	Fresh R01 R02 R05; independent local assessment and server regressions	Bonus not mandatory; JavaScript distinct; known years retained	Canopy optional tools separate; WeDev JavaScript not Java; 2 and 5 year minima recovered	Passed	E1 E4 E5
F01	Result relevance	QA alternate title actions	Live SDET action; 20 fresh source descriptions against persona	Better targeting without false eligibility	Action searches SDET; cohort still 1 strong 5 partial 14 low; no measured relevance uplift	Partially passed	E1 E2 E4
F02	Eligibility	Confirmed facts separate from skills	Four explicit conflicts and unknown controls locally; live disclosure	Country restrictions respected for confirmed user	Policy controls pass; no live synthetic candidate write permitted	Partially passed	E1 E3 E4 E5
F03	Search	Exact filtering and alternate titles	Seven public scenarios; live padded QA empty and SDET	Consistent query effects and clear scope	Padded QA 8; empty query recovery works; exclusion syntax and semantic search deferred	Partially passed	E1 E2 E5
F04	District filter	Preserve office evidence	S3 public replay plus local multi-office tests	Tel Aviv results respect filter	6 results; existing map and district regressions pass; no original defect alleged	Passed	E1 E5
F05	Work style filter	B01 source recovery	Live strict hybrid mobile map; source style controls	Known styles preserved and conflicts unknown	2 strict hybrid results; accessible map links same two jobs	Passed	E1 E2 E5
F06	Country and missing data	Expose known and unknown values	S4 and S4b public replay; local country filters	Strict unknown exclusion and no worldwide inference	2 inclusive remote QA versus 0 strict known country; clear limited-country disclosure	Passed	E1 E2 E5
F07	Salary and seniority filters	Comparable units only	S5; four pay cases; local units and missing controls	Strict ILS/month does not match USD/year or unknown	S5 zero honestly; local comparable and cross-unit tests pass	Passed	E1 E4 E5
F08	Sorting	Label loaded-result scope	Live source to relevance sort after loading more	Stable loaded list; no global ranking implication	99 cards retained; URL sort=relevance; global calibrated catalog ranking absent	Partially passed	E2 E5
F09	Pagination	Snapshot and stale cursor handling	Fresh public S6 two pages; live load more; synthetic cursor tests	No repeated IDs and useful expiry recovery	100 unique public IDs; live 49 to 99 unique cards; local expiry recovery passes	Passed	E1 E2 E5

Audit ID	Original issue	Claimed fix	Test performed	Expected result	Actual result	Verdict	Evidence
F10	Match explanations	Required and preferred quotes and gaps	R03 live false-gap check; exact persona evidence; scoped fixtures	Gaps tied to qualifications; uncertainty explicit	B02 fixed; evidence dimensions and explicit blockers separate; no hiring-probability claim	Passed	E1 E2 E4 E5
F11	Numeric match accuracy	Withhold insufficient evidence	Exact persona calculation and threshold controls	Empirical accuracy against independent truth labels	7/20 numeric scores computed locally; arithmetic controls pass; no empirical calibration dataset	Blocked / not tested	E4 E5
F12	Job descriptions	Safe semantic rendering	Live Canopy tags lists and mobile; local unsafe markup	Readable sections with original semantics	Entity and list repairs work; two source h2 labels remain paragraphs as B03	Partially passed	E2 E3 E5
F13	Data completeness	Recover source facts and unknowns	20 fresh records; compare 14 official APIs	Available facts retained without invented data	17/20 known styles; 3 structured pay plus 2 unstructured; source gaps still present	Partially passed	E1 E3 E4
F14	Freshness disclosure	Separate source and check dates	Live summary versus 26 per-source checks; 20 details	Dates accurately labelled; stale sources disclosed	Card/source dates distinct and stale warning works; summary uses oldest check under last check label N01	Partially passed	E1 E2 Fig 4
F15	Application handoff	Trusted job-specific paths and gate labels	20 original GETs; 14 visible pages; six empty forms; local redirect tests	Correct vacancy path; no false submission claim	20 distinct trusted routes; live apply redirect deliberately uncalled as it logs an event; final Jobicy routes untested	Partially passed	E1 E3 E5
F16	Saved jobs	Reversible persisted state	Reran local save reload unsave and isolation suites	State survives refresh and navigation	Local persistence checks pass; live writes explicitly untested by authorization	Partially passed	E5
F17	Hide and track jobs optional	Company exclusion retained; per-job stages deferred	Inspect live controls and rerun exclusion tests	Per-job hide and self-reported stages if implemented	No per-job hiding or application-stage feature; recommendation deferred not a new regression	Failed	E2 E5
F18	Onboarding and profile	Synthetic CV then preference gate and confirmed criteria	Local first-time workflow and criteria reload suites	Ordered setup; confirmed fields persist	Local workflows pass; hosted first OAuth account and profile writes untested	Partially passed	E5
F19	Resume upload and parsing	PDF DOCX validation and review boundary	Local actual fixture file selection parsing and invalid input tests	Parse valid files and reject invalid inputs before confirmation	Local PDF/DOCX checks pass; live upload not performed; remote scanning unavailable	Partially passed	E5

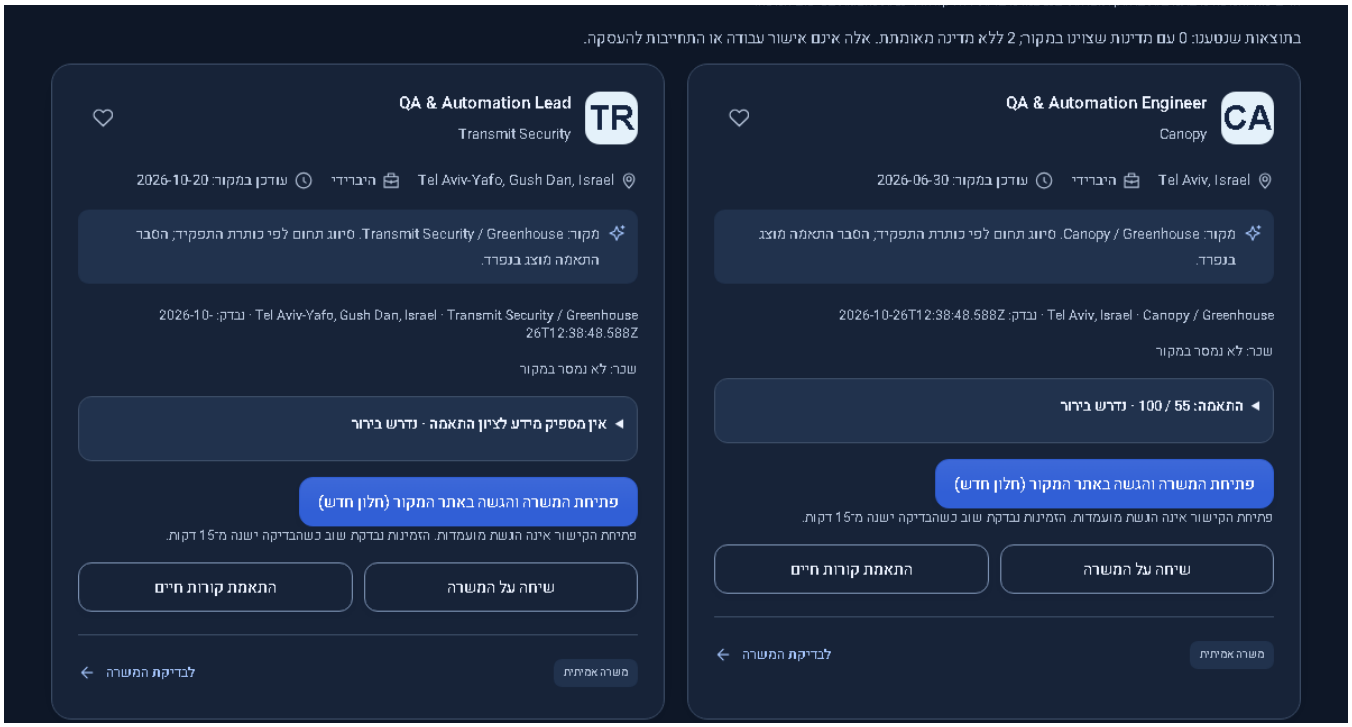
Audit ID	Original issue	Claimed fix	Test performed	Expected result	Actual result	Verdict	Evidence
F20	Mobile layout	Move search above long source panel	Live 390 by 844 plus local themes and device tests	Search visible; no horizontal overflow	Search bottom y509; scroll width 375 below 390; list/map show same two results	Passed	E2 E5 Fig 1
F21	Usability	Chips removal results link remedies	Live style-chip removal Back refresh results jump	State follows controls and history	2 to 3 results after removal; Back and refresh restore 2; results link retains route and query	Passed	E2 E5
F22	Accessibility spot checks	Labels keyboard and structure	Live keyboard skip Enter; mobile link; four local axe scans	Focus main without losing query; semantic document	Skip activation focuses main and preserves URL despite AX href; axe checks pass; B03 heading loss remains; screen reader N/T	Partially passed	E2 E5
F23	Performance benchmark	Bounded metadata residency	Built current code; observed request times; no field benchmark	Repeatable field latency memory and loading metrics	Public scenario GETs about 0.7 to 1.0s; live refresh reached stale fallback; no controlled field benchmark	Blocked / not tested	E1 E2 E5
F24	Empty and recovery states	Specific restrictions and retry	Live nonsense query then QA; real stale source; local cursor faults	Clear empty state and recovery without fabricated listings	0 to 8 on query recovery; stale-source warning with retained cards; isolated recovery tests pass	Passed	E2 E5
F25	Errors and outages	Fail truthfully and recover	Local provider network storage abort failure suites; actual stale feed	Explicit failure no substitute fake jobs; retry works	Isolated failures pass and live stale copy disclosed; no forced production outage or quota exhaustion	Partially passed	E2 E5
F26	AI assistant	Remain unavailable without provider and consent	Local missing provider and mocked recovery tests	Only configured authorized real model grounded output	Real model quality not tested; provider consent and cost authorization required	Blocked / not tested	E5 E7

Before and after discovery evidence

Figure 1 relates to B01 and F20. The before capture below is the supplied original audit page 4, not a newly reproduced failure. The after card crop is a fresh isolated browser regression execution using public vacancy fixtures, with their historical source timestamps preserved. Independent live strict-hybrid and map checks also returned the same two jobs. The live mobile screenshot excludes the account header.



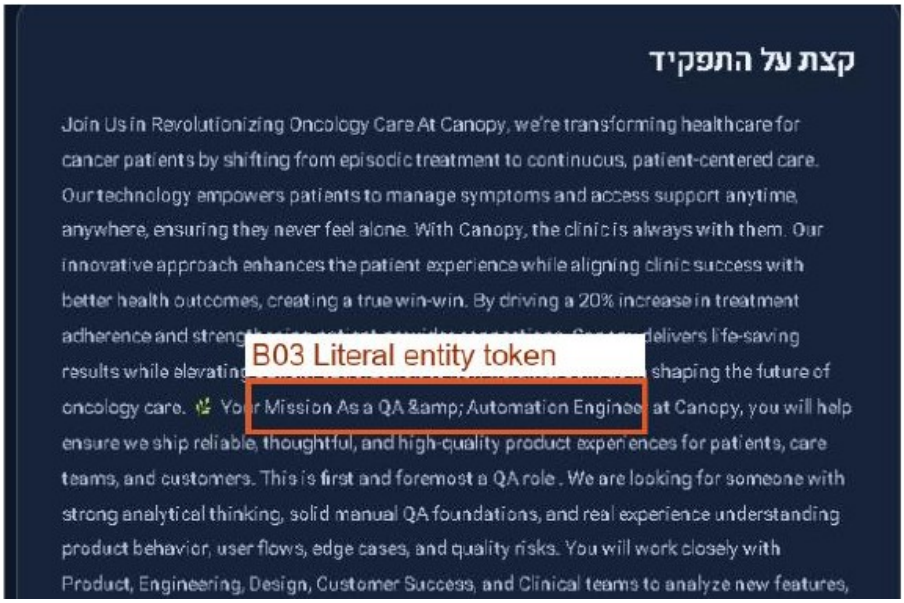
Original audit screenshot page 4 — Canopy work arrangement labelled unknown. B01.



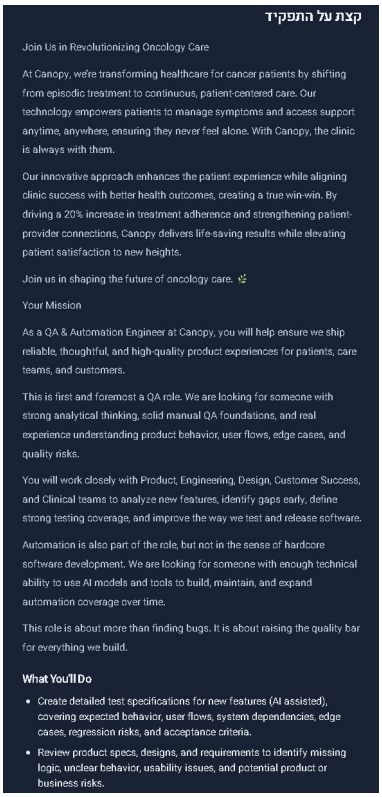
After — current local UI regression, two hybrid source vacancies. Synthetic profile; not a live account mutation.

Description verification evidence

Figure 2 for B03. The before screenshot is the original audit's annotated description from page 14. The after crop contains only public live job text. Entities are decoded and lists are restored. The employer source marks both the opening title and Your Mission as h2; the current Niva DOM renders these as p. What You'll Do, What You'll Bring, Bonus Points and Benefits are h3, with 27 list items. This is an incomplete repair, not a new regression.



Original audit screenshot page 14 — encoded amp token and flattened description.



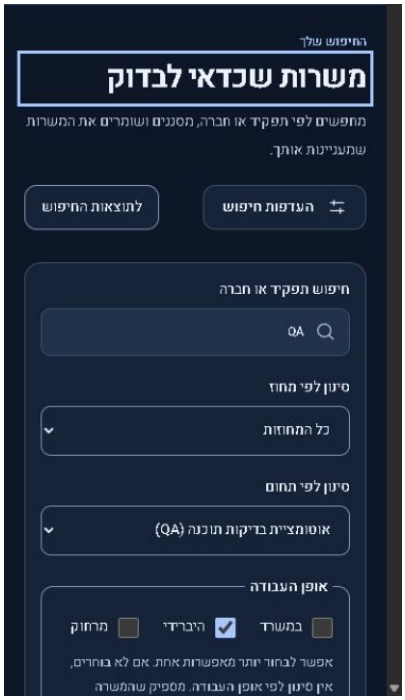
After — current live Canopy description. Your Mission remains paragraph text; personal account areas excluded.

Pay and mobile verification evidence

Figure 3 for B04 and F20. The salary capture is a fresh local isolated regression using the public R09 fixture. Employer-estimated annual USD pay is disclosed rather than converted to the persona’s ILS/month target. The mobile capture is current live390×844, cropped below the account header; search bottom was y509 and no horizontal overflow was observed.



After — isolated R09 salary summary with employer estimate.



After — live mobile search first. Header removed to redact account identity. F20.

New issue and regression checks

N01 Oldest check labelled as last check

Severity Low P2. Previously undocumented; the version that introduced it is not established. Steps: open jobs with scope=all, expand source coverage, expand source check dates, compare the summary with per-source dates. Expected: a last-check label identifies the latest successful observation, or the UI explicitly labels the oldest source date and shows the range. Actual: summary shows 25 October 18:27:09 UTC while many sources show 26 October 16:57 UTC. Code corroboration: server/discovery.ts returns the ascending minimum checked_at as fetchedAt; jobs.html labels it last source check. This is misleading presentation, not evidence a newer listing is stale or that retention failed.



Figure 4 — current live summary and newer per-source timestamps. N01. Public source data only.

Related behavior tested

Style-chip removal changed 2 to 3 results; Back and refresh restored the original 2. The results jump kept the route and query. Padded QA returned 8, a nonsense query returned 0, and QA recovery returned 8. SDET sets an explicit alternate title and automation category. Loaded-result sorting retained 99 cards and changed the URL sort parameter. Public pagination had 100 unique IDs. The live source-failure warning appeared naturally while older listings remained visible.

A suspected global skip-link defect was not confirmed: its accessibility URL looked like /#main, but Enter focused the current main and preserved the jobs route and query. It is not registered as a failure. Full local suite fault cases covered aborted requests, stale cursors, unavailable providers, storage failures and recovery. No production outage or rate-limit failure was injected. No new high-severity regression was established; B03 remains incomplete and N01 is newly documented.

Updated feature scores

Scores are analyst judgments using the audit's1–10 convention, not test-pass percentages. Live scores use current public behavior supported by isolated controls. N/T is used where live account writes, calibration or field behavior were not tested; a successful local test is not silently promoted to a live score. Original N/T values remain distinct from zero.

ID	Feature	Original /10	Verified live /10
F01	Result relevance	6	6
F02	Eligibility	5	7
F03	Search	8	8
F04	District filter	8	8
F05	Work style filter	5	8
F06	Country and missing data	7	8
F07	Salary and seniority filters	6	7
F08	Sorting	6	6
F09	Pagination	8	8
F10	Match explanations	4	8
F11	Numeric match accuracy	N/T	N/T
F12	Job descriptions	5	7
F13	Data completeness	4	7
F14	Freshness disclosure	7	6
F15	Application handoff	8	8
F16	Saved jobs	8	N/T
F17	Hide and track jobs optional	N/T	N/T
F18	Onboarding and profile	N/T	N/T
F19	Resume upload and parsing	N/T	N/T
F20	Mobile layout	8	8
F21	Usability	6	8
F22	Accessibility spot checks	7	7
F23	Performance benchmark	N/T	N/T
F24	Empty and recovery states	8	8
F25	Errors and outages	N/T	N/T
F26	AI assistant	N/T	N/T

Local-only coverage remains strong for F16 saved jobs, F18 setup, F19 PDF/DOCX parsing and F25 faults, but those live paths are N/T. F11 arithmetic is exercised while empirical score accuracy remains N/T. F17 is absent optional functionality. Real AI and field performance remain N/T.

Weighted original rubric

Dimension	Weight	Original	Retest
Relevance	25%	6	6
Eligibility	15%	5	7
Data quality	20%	4	7
Application handoff	15%	8	8
Search	10%	8	8
Usability	10%	6	8
Responsive accessibility	5%	7	7
Weighted total	100%	6.0	7.1

7.1 is provisional within the bounded public journey, not a complete production-readiness score. Handoff 8 reflects job-specific originals and local redirect controls, with the live tracking redirect untested. The fix report's7.3 local provisional score is not directly comparable: this retest reduces data quality for remaining source gaps/heading loss and separately flags freshness disclosure. No relevance uplift was demonstrated.

Supporting records and code references

Evidence is retained under .test-data/verify41: public.json, browser-sources.json, assessments.json, verification.json, frontend-tests.json, server-tests.json, browser-tests.json and execution logs. These are local audit artifacts; they are not production telemetry or application patches. The new verification source and report were authored independently from earlier conclusions.

Live Niva jobs

Reviewed deployed code

Source facts and extraction

Assessment and uncertainty policy

Description component with incomplete headings

Executed checks and remaining gates

Check	Independent execution result
Frontend full suite	59 passed /0 failed; fresh run
Server full suite	153 passed /0 failed; fresh run
Chromium full browser suite	218 passed /0 failed /0 skipped /0 flaky; desktop and mobile; four audit theme/device axe checks
Independent original persona	28 assertions passed on 20 fresh public records; explicit conflicts and exact URLs checked
Worker and source provenance	1/1 each passed; local runtime/contract evidence
Production build	Passed; HTML policy and Worker build completed
Type checks	Server, browser tests and frontend tests passed
Local and hosted boundary	Local SQLite/D1 contract only; not proof of hosted D1 isolation or cleanup

Prioritized acceptance criteria

P1 S02 — assign an accountable operator and approved alert destination; observe two hosted cleanup cycles with expired-row deletion, denial and recovery evidence; restore a disposable authorized account and verify identity isolation. Until then operational release sign-off remains blocked.

P1 F15 F16 F18 F19 — use an authorized staging account to finish first-time identity setup, real PDF/DOCX parsing and confirmation, preferences, save/unsave across refresh and separate sessions; test hosted apply redirect logging and exact final vacancy path without submitting an application. All live write checks remain untested here by instruction.

P2 B03 F12 — preserve source heading boundaries rather than a vocabulary-only heuristic. Recheck Canopy opening heading and Your Mission, all four original affected records, original order, safe HTML removal and 390px reflow; every source heading should retain a semantic heading role.

P2 N01 F14 — label minimum source date as oldest check, or return separately labelled newest/oldest dates. With a two-source fixture having distinct dates, both labels must match; retain stale/failure details without claiming employer confirmation.

P2 F01 F08 F11 — expand a labelled representative QA cohort and publish relevance/eligibility/score calibration separately. Do not interpret more catalog inventory as better matches. Include alternate titles and known negative controls before changing global ranking.

P3 F17 — design reversible per-job hiding and explicitly self-reported stages if prioritized. Opening a link must never mark applied. This is an optional outstanding recommendation, not a new blocker.

Limitations and protected data

No applications, employer contacts, purchases, paid activation, deployment or application-source edits were made. Private owner CV excerpts, identity and numeric assessments are excluded from delivered figures. No physical phone, spoken Hebrew screen reader, browser-engine cross-platform matrix, controlled field benchmark, full-catalog semantic duplicate review, Jobicy final employer gate or real AI-quality evaluation was completed. Rendered document checks certify document layout only, not software behavior.